

Question 4

A. Load the /public/bmort/R/heart.csv data set into a data frame. Are there any missing values? Perform any necessary data imputation on the data set.

In [15]:

```
## Loading the dataset
heart <- read.csv('/public/bmort/R/heart.csv')
head(heart,10)
```

A data.frame: 10 × 14

age	sex	pain	bp	chol	sugar	ecg	rate	angina	stv	sts	mvn	thal	c
<int>	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<int>	<int>
63	1	1	145	233	1	2	150	0	2.3	3	0	6	0
67	1	4	160	286	0	2	108	1	1.5	2	3	3	1
67	1	4	120	229	0	2	129	1	2.6	2	2	7	1
37	1	3	130	250	0	0	187	0	3.5	3	0	3	0
41	0	2	130	204	0	2	172	0	1.4	1	0	3	0
56	1	2	120	236	0	0	178	0	0.8	1	0	3	0
62	0	4	140	268	0	2	160	0	3.6	3	2	3	1
57	0	4	120	354	0	0	163	1	0.6	1	0	3	0
63	1	4	130	254	0	2	147	0	1.4	2	1	7	1
53	1	4	140	203	1	2	155	1	3.1	3	0	7	1

In [16]:

```
## Finding missing values
## Locating the missing values
which(is.na(heart), arr.ind = TRUE)
```

A matrix:

0 × 2 of

type int

row	col
-----	-----

It can be seen that this data has no missing values.

In [17]:

```
### Converting the disease column in to a factor
heart$disease <- as.factor(heart$disease)
```

B. Produce a table of summary statistics on the data set. How do the ranges of the values in the columns compare? Does each column of data have similar magnitudes and ranges? Are there any outliers?

In [18]:

```
## summary statistics
summary(heart)
```

age	sex	pain	bp	chol
Min. :29.00	Min. :0.00	Min. :1.000	Min. : 94.0	Min. :12
6.0				
1st Qu.:48.00	1st Qu.:0.00	1st Qu.:3.000	1st Qu.:120.0	1st Qu.:21
1.0				
Median :56.00	Median :1.00	Median :3.000	Median :130.0	Median :24
1.5				
Mean :54.48	Mean :0.68	Mean :3.153	Mean :131.6	Mean :24
6.9				
3rd Qu.:61.00	3rd Qu.:1.00	3rd Qu.:4.000	3rd Qu.:140.0	3rd Qu.:27
5.2				
Max. :77.00	Max. :1.00	Max. :4.000	Max. :200.0	Max. :56
4.0				

sugar	ecg	rate	angina
Min. :0.0000	Min. :0.0000	Min. : 71.0	Min. :0.0000
1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:133.8	1st Qu.:0.0000
Median :0.0000	Median :0.5000	Median :153.0	Median :0.0000
Mean :0.1467	Mean :0.9867	Mean :149.7	Mean :0.3267
3rd Qu.:0.0000	3rd Qu.:2.0000	3rd Qu.:166.0	3rd Qu.:1.0000
Max. :1.0000	Max. :2.0000	Max. :202.0	Max. :1.0000

stv	sts	mvn	thal	disease
Min. :0.00	Min. :1.000	Min. :0.00	Min. :3.000	0:162
1st Qu.:0.00	1st Qu.:1.000	1st Qu.:0.00	1st Qu.:3.000	1:138
Median :0.80	Median :2.000	Median :0.00	Median :3.000	
Mean :1.05	Mean :1.603	Mean :0.67	Mean :4.727	
3rd Qu.:1.60	3rd Qu.:2.000	3rd Qu.:1.00	3rd Qu.:7.000	
Max. :6.20	Max. :3.000	Max. :3.00	Max. :7.000	

Columns like sex,sugar, angina and ecg are said to have nominal outcomes like(0,1) and (0,2). Pain on the otherhand shows a ordinal outcomes and these variables are said to be categorical data. Same can be said for other qualitative variables in out data.

bp, chol, considering the 3rd quartile and maximum depicts the existence of upper outliers. rate , considering the 1st quartile and minimum shows a lower outlier.

Age has no outliers.

C. Partition the heart data set so that 80% will be used for training and 20% will be used for testing your machine learning model.

In [24]:

```
install.packages('caret')
install.packages('ggplot2')
install.packages('lattice')
library(caret)
library(ggplot2)
library(lattice)
```

In [20]:

```
### Splitting the dataset
sp_data <- createDataPartition(y = heart$disease, p = 0.8, list = FALSE)
# sp_data
```

In [21]:

```
## The training and testing data
tr_data <- heart[sp_data,]
te_data <- heart[-sp_data,]
```

D. Using logistic regression as provided by the Caret library in R, develop a model to predict heart disease diagnosis based on the 13 features provided in the data set for each patient.

In [23]:

```
## Fitting a logistic regression model
log_model <- train(disease~., data = tr_data, method = 'glm', family = 'binomial')
log_model
```

Generalized Linear Model

241 samples
13 predictor
2 classes: '0', '1'

No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 241, 241, 241, 241, 241, 241, ...
Resampling results:

Accuracy	Kappa
0.8249149	0.6463876

E. Generate a confusion matrix using the data from your test set to show the accuracy of the model.

In [29]:

```
## predicting the test data
pred_test1 <- predict(object = log_model, newdata = te_data)
pred_test1
```

```
0 1 1 0 0 1 1 0 1 1 0 0 0 1 0 0 0 1 0 0 1 1 1 1
1 0 0 0 0 1 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1
0 1 0 0 0 0 0 1 1 1 0
```

► Levels:

In [30]:

```
confusionMatrix(table(pred_test1, te_data$disease))
```

Confusion Matrix and Statistics

```
pred_test1  0  1
            0 29  9
            1  3 18
```

```
          Accuracy : 0.7966
          95% CI   : (0.6717, 0.8902)
    No Information Rate : 0.5424
    P-Value [Acc > NIR] : 4.294e-05
```

```
          Kappa : 0.583
```

```
McNemar's Test P-Value : 0.1489
```

```
          Sensitivity : 0.9062
          Specificity : 0.6667
    Pos Pred Value   : 0.7632
    Neg Pred Value   : 0.8571
          Prevalence : 0.5424
    Detection Rate   : 0.4915
    Detection Prevalence : 0.6441
    Balanced Accuracy : 0.7865
```

```
'Positive' Class : 0
```

F. Write a few sentences providing commentary on the accuracy of the model. What percent are false positives? What percent are false negatives?

The accuracy of the log_model has a **79.66%** classification rate. From the confusion matrix above, we can see that the model correctly classified 29 non diseased patients as no disease whiles it wrongly classified 9 not diseased patients as disease. The model also wrongly classified 3 diseased patients as no disease whiles it correctly classified 18 diseased patients as disease.

In [32]:

```
conf_table <- table(pred_test1, te_data$disease)
conf_table
```

```
pred_test1  0  1
            0 29  9
            1  3 18
```

In [40]:

```
false_pos <- conf_table[2,1]
per_fp <- round((false_pos/sum(conf_table[2,]))*100,2)
per_fp
```

14.29

The percent of false positives are **14.29%**.

In [41]:

```
false_neg <- conf_table[1,2]
per_np <- round((false_neg/sum(conf_table[1,]))*100,2)
per_np
```

23.68

The percent of false negatives are **23.68%**.